Node Structure Visualization

Click on any node in the tree to jump to its detailed information.
Total Nodes: 18 | Best Validation Score: 0.9176 (Node 13)

Node Tree:

Table of Contents:

Node ID	Stage	Status	Tool
-1	root		
0	evolve	✗	autogluon.multimodal
1	evolve	✗	machine learning
2	evolve	✗	autogluon.tabular
3	debug	✗	autogluon.multimodal
4	debug	✗	machine learning
5	debug	✗	autogluon.multimodal
6	debug	✗	machine learning
7	debug	✗	autogluon.tabular
8	debug	✗	autogluon.multimodal
9	debug	✗	autogluon.multimodal
10	debug	✗	machine learning
11	debug	✗	autogluon.tabular
12	debug	✗	autogluon.multimodal
13	debug	✓	machine learning
14	debug	✗	autogluon.tabular
15	debug	✗	autogluon.tabular
16	debug	✗	autogluon.tabular

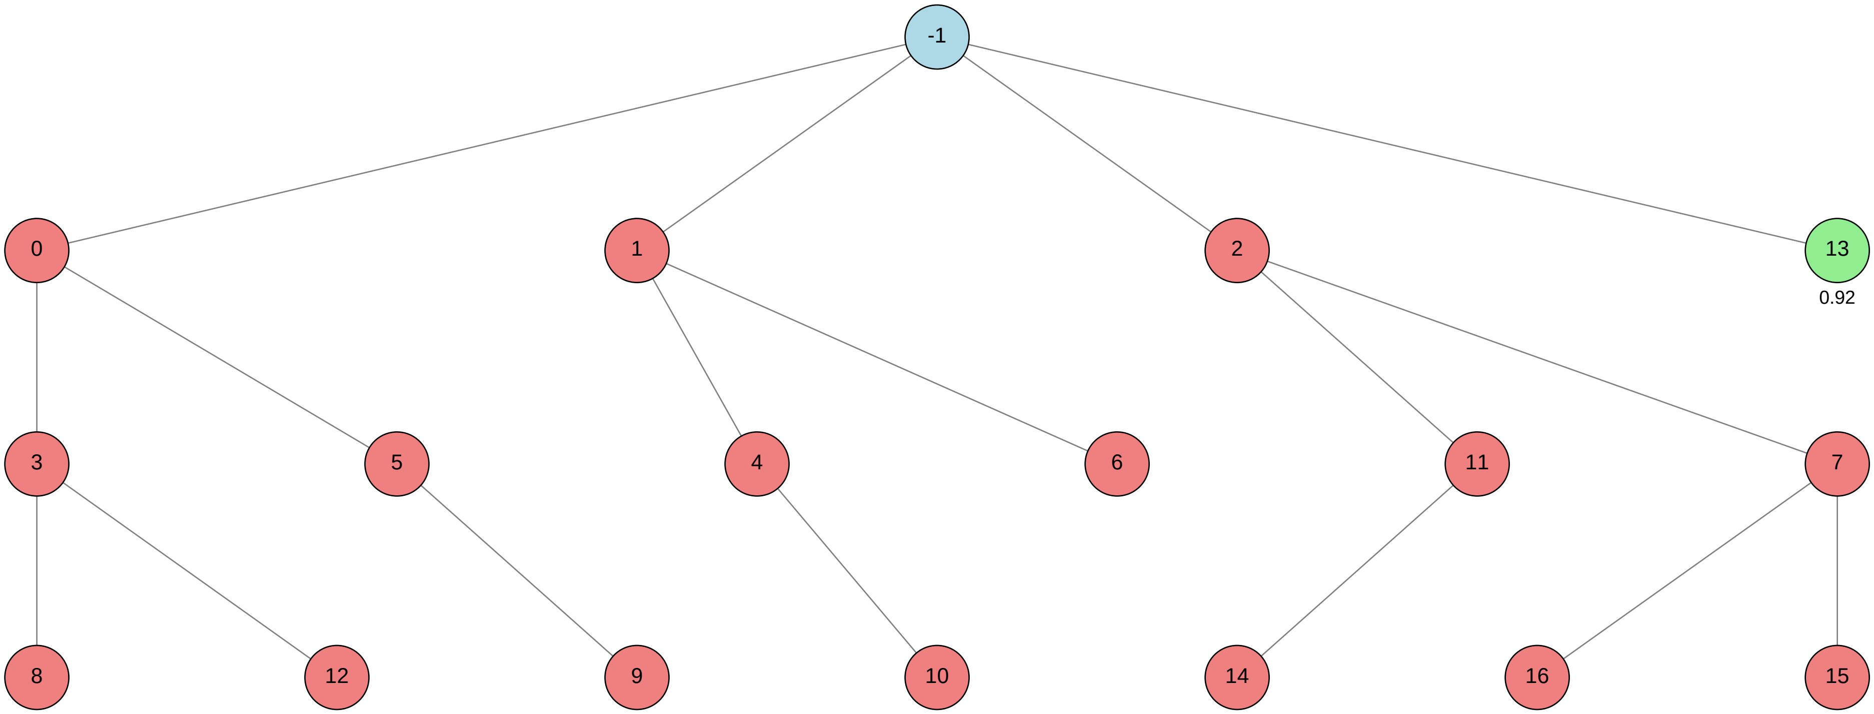

Status	Color
Success	Green
Failure	Red
Neutral	Blue

Node Details:

Node -1 (root)

Property	Value
uct_value	0.5882
ctime	2026-05-21 15:58:24
debug_attempts	0
depth	0
execution_time	0.0
expected_child_count	0
failure_visits	16
id	-1
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
num_children	4
prev_tutorial_prompt	None
stage	root
time_step	-1
tool_used	
tools_available	[]
unvalidated_visits	0
validated_reward	0.9176
validated_visits	1
validation_score	None
visits	17
parent_id	None
child_ids	[0, 1, 2, 13]

Node 0 (evolve)

Property	Value
uct_value	0.6383
ctime	2026-05-21 15:59:41
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.3
exploration	0.9968262051089475
failure_penalty	-0.3
failure_visits	6
id	0
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	3
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	0
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	6
parent_id	-1
child_ids	[3, 5]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_0/conda_env
Resolved 235 packages in 3.55s
Uninstalled 1 package in 1.68s
Installed 233 packages in 43.96s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0

+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7

+ idna==3.1... [truncated, 10537 chars total]

Error Analysis:

ERROR_SUMMARY: The training process was terminated because it exceeded the allowed wall-clock time limit of 3600 seconds (1 hour). As a result, the model did not complete training, and subsequent steps such as prediction and output file generation were not executed, leading to missing required output files.

SUGGESTED_FIX:
Reduce the training time by setting the `time_limit` parameter in `predictor.fit()` to a value less than 3600 seconds (e.g., 3300–3500 seconds to allow time for data loading and inference). Additionally, consider using a faster preset (such as `"medium_quality_faster_train"`), reducing the number of epochs, or using a smaller model backbone to ensure the entire workflow (data loading, training, and prediction) completes within the time constraint.

Node 1 (evolve)

Property	Value
uct_value	0.9400
ctime	2026-05-21 17:03:24
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.25
exploration	0.849219727985642
failure_penalty	-0.25
failure_visits	4
id	1
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	2
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	1
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	4
parent_id	-1
child_ids	[4, 6]

Error Message:

```
stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 29-basic_prompt/node_1/conda_env
Resolved 79 packages in 366ms
Uninstalled 1 package in 1.72s
Installed 78 packages in 33.56s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.0
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0
+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.0
+ idna==3.15
+ jinja2==3.1.6
```

+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocess==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparse==12.6.3.3
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.13.4
+ pydantic-core==2.46.4
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2
+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 4418 chars total]

Error Analysis:

ERROR_SUMMARY: The error occurs because the 'verbose' argument is not supported in the version of PyTorch installed in your environment for the `ReduceLROnPlateau` learning rate scheduler. Passing this unsupported keyword argument causes a `TypeError` during initialization.

SUGGESTED_FIX: Remove the `verbose=True` argument from the call to `torch.optim.lr_scheduler.ReduceLROnPlateau` in your script. After removing this argument, rerun your script to ensure compatibility with your current PyTorch version.

Node 2 (evolve)

Property	Value
uct_value	0.6383
ctime	2026-05-21 17:08:49
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.3
exploration	0.5264742539237474
failure_penalty	-0.3
failure_visits	6
id	2
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	3
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	2
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	6
parent_id	-1
child_ids	[11, 7]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_2/conda_env
Resolved 239 packages in 1.40s
Uninstalled 1 package in 1.99s
Installed 237 packages in 46.10s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0

+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2

+ hyperopt==... [truncated, 41834 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes, resulting in OSError: "validate_socket_filename failed: AF_UNIX path length cannot exceed 107 bytes". This caused all model training attempts to fail, leading to a RuntimeError: "No models were trained successfully during fit()".

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output, especially the base working directory and any subdirectories created by your script or environment. Move your project and scratch directories to a location with a much shorter absolute path (e.g., /tmp/ml or /scratch/ml), and update all relevant path variables in your scripts and environment setup accordingly. Then rerun the training; this should allow Ray and AutoGluon to create socket files with paths under the 107-byte limit, resolving the error.

Node 3 (debug)

Property	Value
uct_value	0.9261
ctime	2026-05-21 17:14:54
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.2684446626418298
failure_penalty	-0.0
failure_visits	3
id	3
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	2
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementatio
stage	debug
time_step	3
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	3
parent_id	0
child_ids	[8, 12]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_3/conda_env
Resolved 235 packages in 1.26s
Uninstalled 1 package in 2.07s
Installed 233 packages in 47.04s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10

+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0

+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.1... [truncated, 9368 chars total]

Error Analysis:

ERROR_SUMMARY: The script failed with `ValueError: Unknown preset type: medium_quality_faster_train` during the call to `predictor.fit()`. This indicates that the string `"medium_quality_faster_train"` is not a recognized value for the `presets` argument in the installed version of AutoGluon MultiModal, likely because this preset does not exist or has been renamed/removed in your package version.

SUGGESTED_FIX:
Check the available preset names for your installed version of AutoGluon MultiModal (e.g., by consulting the [official documentation](https://auto.gluon.ai/stable/tutorials/multimodal/index.html) or inspecting `autogluon.multimodal.utils.presets.AUTOMM_PRESETS`). Replace `"medium_quality_faster_train"` with a valid preset such as `"medium_quality_faster_inference"`, `"best_quality"`, `"medium_quality"`, or simply remove the `presets` argument to use the default. Then rerun the script.

Node 4 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-21 17:21:49
debug_attempts	2
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.1772322201769845
failure_penalty	-0.0
failure_visits	2
id	4
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure yo # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have ther
stage	debug
time_step	4
tool_used	machine learning

tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	2
parent_id	1
child_ids	[10]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt` does not contain any dependencies
Using Python 3.11.15 environment at: 29-basic_prompt/node_4/conda_env
Resolved 79 packages in 876ms
Uninstalled 1 package in 1.82s
Installed 78 packages in 42.03s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.0
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0

+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.0
+ idna==3.15
+ jinja2==3.1.6
+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocess==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparse==12.6.3.3
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.13.4
+ pydantic-core==2.46.4
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2

+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 89101 chars total]

Error Analysis:

ERROR_SUMMARY: The script was terminated because it exceeded the 1-hour wall-clock time limit imposed by the environment. This caused the training process to be killed before all epochs completed, resulting in incomplete model training and missing final outputs. There were no explicit code errors, but the workflow could not finish due to insufficient time for the full training and inference pipeline.

SUGGESTED_FIX:
Reduce the overall runtime by decreasing the number of training epochs, lowering the batch size, or using a smaller model architecture. Additionally, consider reducing the number of DataLoader workers, optimizing data loading, or performing early stopping based on validation performance to ensure the script completes within the 1-hour time limit.

Node 5 (debug)

Property	Value
uct_value	1.3384
ctime	2026-05-21 18:24:55
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.664857771836881
failure_penalty	-0.0
failure_visits	2
id	5
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementatio
stage	debug
time_step	5
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	2
parent_id	0
child_ids	[9]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_5/conda_env
Resolved 235 packages in 1.18s
Uninstalled 1 package in 1.74s
Installed 233 packages in 58.97s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10

+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0

+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.1... [truncated, 17360 chars total]

Error Analysis:

ERROR_SUMMARY: The error is an `OSError: [Errno 16] Device or resource busy` raised by Python's `shutil.rmtree` during the cleanup of temporary directories used by multiprocessing. This occurs because files with names like `.nfs*` are still open or locked by another process or thread (often on NFS-mounted filesystems), preventing their deletion at process exit. This is a known issue with NFS and does not affect the main training, prediction, or output generation, as all validation checks passed and the model was successfully created and saved.

SUGGESTED_FIX:
You can safely ignore this error if your main outputs (model and predictions) are correct, as it only affects temporary file cleanup. To reduce or eliminate these warnings, ensure all file handles are properly closed before process exit, avoid using NFS for temp directories if possible, or set the `TMPDIR` environment variable to a local disk. If the warnings are disruptive, you may also suppress them by catching and ignoring `OSError` during cleanup, but this is not strictly necessary for correct results.

Node 6 (debug)

Property	Value
uct_value	1.6649
ctime	2026-05-21 18:54:23
debug_attempts	2
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.664857771836881
failure_penalty	-0.0
failure_visits	1
id	6
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	0
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure yo # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have ther
stage	debug
time_step	6
tool_used	machine learning

tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	1
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt` does not contain any dependencies
Using Python 3.11.15 environment at: 29-basic_prompt/node_6/conda_env
Resolved 79 packages in 1.01s
Uninstalled 1 package in 2.30s
Installed 78 packages in 40.11s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.0
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0

+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.0
+ idna==3.15
+ jinja2==3.1.6
+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocess==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparse==12.6.3.3
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.13.4
+ pydantic-core==2.46.4
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2

+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 88935 chars total]

Error Analysis:

ERROR_SUMMARY: The script was terminated before completion because it exceeded the 1-hour wall-clock time limit imposed on the process. This prevented the training from finishing all epochs and blocked subsequent steps such as final model saving, test set prediction, and output validation. There were no explicit code errors, but the workflow could not complete due to insufficient time for the combined data loading, training, and inference.

SUGGESTED_FIX:
Reduce the total training time by lowering the number of training epochs, decreasing the model size (e.g., use a smaller or more efficient architecture), reducing the batch size, or limiting the number of DataLoader workers. Additionally, consider saving intermediate results and performing test set inference as soon as a reasonable model checkpoint is available, rather than waiting for all epochs to finish, to ensure that predictions and outputs are generated before the time limit is reached.

Node 7 (debug)

Property	Value
uct_value	0.9261
ctime	2026-05-21 19:56:51
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.2684446626418298
failure_penalty	-0.0
failure_visits	3
id	7
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	2
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinde
stage	debug
time_step	7
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	3
parent_id	2
child_ids	[16, 15]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_7/conda_env
Resolved 239 packages in 1.31s
Uninstalled 1 package in 1.99s
Installed 237 packages in 1m 09s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10

+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1

+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==... [truncated, 41815 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes, resulting in repeated OSError exceptions: "AF_UNIX path length cannot exceed 107 bytes". This caused all model training attempts to fail, leading to a RuntimeError: "No models were trained successfully during fit()".

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output, especially the scratch and output directories, so that the full UNIX socket paths remain under 107 bytes. Move your working directory (including data, output, and conda environment) to a location with a much shorter absolute path (e.g., /tmp/ml or /scratch/ml), and re-run the script to avoid exceeding the UNIX socket path length limit.

Node 8 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-21 20:03:49
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	8
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementatio
stage	debug
time_step	8
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	3
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_8/conda_env
Resolved 235 packages in 1.10s
Uninstalled 1 package in 1.72s
Installed 233 packages in 51.78s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0

```
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.1... [truncated, 17358 chars total]
```

Error Analysis:

ERROR_SUMMARY: The main code executed successfully, but during cleanup, multiple `OSError: [Errno 16] Device or resource busy` errors occurred when attempting to delete temporary files (with `.nfs*` filenames) on an NFS-mounted filesystem. This happens when files are still open by a process while being deleted, causing NFS to rename them instead of removing them, and subsequent deletion attempts fail because the files are still in use.

SUGGESTED_FIX:
These errors are benign and do not affect model training or prediction results. To suppress them, ensure all file handles are properly closed before process exit, and avoid manual deletion of files still in use. If possible, run the script on a local filesystem or consult your system administrator to adjust NFS settings or cleanup policies for `.nfs*` files.

Node 9 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-21 20:31:44
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	9
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementatio
stage	debug
time_step	9
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	5
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_9/conda_env
Resolved 235 packages in 2.88s
Uninstalled 1 package in 1.75s
Installed 233 packages in 53.90s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.1... [truncated, 17372 chars total]

Error Analysis:

ERROR_SUMMARY: The main error is an `OSError: [Errno 16] Device or resource busy` raised during the cleanup of temporary directories by Python's multiprocessing module. This occurs when the script attempts to delete files (with names like `.nfs*`) that are still open or in use, typically on NFS-mounted filesystems, after the main training and prediction tasks have completed successfully.

SUGGESTED_FIX:
This error is generally benign and does not affect the correctness of your model or predictions. To suppress these warnings, ensure all file handles are properly closed before script exit, and avoid manual deletion of files that may still be in use. If possible, run your workflow on a local filesystem instead of NFS, or consult your system administrator to adjust NFS settings for handling deleted-but-open files.

Node 10 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-21 20:59:02
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	10
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have ther
stage	debug
time_step	10
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0

validation_score	None
visits	1
parent_id	4
child_ids	[]

Error Message:

```
stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 29-basic_prompt/node_10/conda_env
Resolved 79 packages in 1.17s
Prepared 1 package in 900ms
Uninstalled 1 package in 1.73s
Installed 78 packages in 37.20s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.0
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0
+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.1
```

```
+ idna==3.15
+ jinja2==3.1.6
+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocess==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparse==12.6.3.3
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.13.4
+ pydantic-core==2.46.4
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2
+ rich==15.0... [truncated, 89641 chars total]
```

Error Analysis:

ERROR_SUMMARY: The process was terminated because it exceeded the 1-hour wall-clock time limit before completing all training epochs and subsequent prediction/saving steps. This was likely due to a combination of a large dataset, a relatively slow model (EfficientNet-B0), high batch size, and excessive DataLoader worker count, possibly exacerbated by an outdated NVIDIA GPU driver causing suboptimal CUDA performance.

SUGGESTED_FIX:
Reduce the number of training epochs (e.g., from 8 to 3–4), lower the DataLoader batch size and number of workers (e.g., batch_size=32, num_workers=4–8), and consider using a faster or smaller model architecture (e.g., resnet18 or mobilenet_v2). Additionally, update the NVIDIA GPU driver to the latest version to ensure optimal CUDA performance, and monitor training time per epoch to ensure the script completes well within the 1-hour limit.

Node 11 (debug)

Property	Value
uct_value	1.3384
ctime	2026-05-21 22:01:26
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.482079962591042
failure_penalty	-0.0
failure_visits	2
id	11
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinde
stage	debug
time_step	11
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	2
parent_id	2
child_ids	[14]

Error Message:

```
stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_11/conda_env
Resolved 239 packages in 2.95s
Uninstalled 1 package in 1.78s
Installed 237 packages in 51.77s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
```

+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1

+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt=... [truncated, 41799 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend attempted to create UNIX domain socket files with paths exceeding the 107-byte limit imposed by the AF_UNIX socket protocol. This resulted in repeated OSError exceptions (validate_socket_filename failed), causing all model training attempts to be skipped and ultimately raising a RuntimeError indicating that no models were successfully trained.

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output, especially the scratch and output directories, so that the full socket file paths generated by Ray remain under 107 characters. Move your working directory (e.g., /sc-scratch/sc-scratch-ikim-guidlight/be-fair/mlzero/29-basic_prompt/node_11/) to a location with a much shorter absolute path (such as /tmp/agrun or /scratch/agrun), update all relevant path variables in your scripts accordingly, and rerun the pipeline.

Node 12 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-21 22:07:55
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	12
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementatio
stage	debug
time_step	12
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	3
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_12/conda_env
Resolved 235 packages in 1.61s
Uninstalled 1 package in 1.81s
Installed 233 packages in 1m 16s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.... [truncated, 18060 chars total]

Error Analysis:

ERROR_SUMMARY: The main error is an AttributeError: 'MultiModalPredictor' object has no attribute 'leaderboard', which occurs because the leaderboard() method is not available for MultiModalPredictor in the installed version of AutoGluon. Additionally, repeated OSError: [Errno 16] Device or resource busy errors related to .nfs* files are seen during temporary directory cleanup, which are common on NFS filesystems when files are still open by a process.

SUGGESTED_FIX:
Remove or comment out the call to predictor.leaderboard(), as this method is not supported for MultiModalPredictor; use predictor.evaluate() instead to obtain validation metrics if needed. The OSError messages about .nfs* files can generally be ignored unless they cause disk space/resource issues, but if they are problematic, ensure all file handles are properly closed before cleanup or consult your system administrator for NFS tuning.

Node 13 (debug)

Property	Value
uct_value	3.3801
ctime	2026-05-21 22:36:02
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	0
id	13
is_debug_successful	✗
is_leaf	✓
is_successful	✓
is_terminal	✓
num_children	0
prev_tutorial_prompt	None
stage	debug
time_step	13
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.9176
validated_visits	1
validation_score	0.9176
visits	1
parent_id	-1
child_ids	[]

Node 14 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-21 23:27:52
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	14
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial explains GPU acceleration in AutoGluon, covering basic GPU allocation with `num_gpus
stage	debug
time_step	14
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	11
child_ids	[]

Error Message:

```
stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_14/conda_env
Resolved 239 packages in 1.36s
Uninstalled 1 package in 1.73s
Installed 237 packages in 46.11s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
```

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt=... [truncated, 41852 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes. This resulted in repeated OSError exceptions ("AF_UNIX path length cannot exceed 107 bytes"), causing all model training attempts to fail and ultimately raising a RuntimeError: "No models were trained successfully during fit()".

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output—specifically, move your working, scratch, and output directories to locations with much shorter absolute paths (e.g., /tmp/ag or /scratch/ag). This will ensure that the UNIX socket paths created by Ray remain under the 107-byte limit and allow AutoGluon to train models successfully.

Node 15 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-21 23:33:37
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	15
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinde
stage	debug
time_step	15
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	7
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_15/conda_env
Resolved 239 packages in 1.13s
Uninstalled 1 package in 1.69s
Installed 237 packages in 1m 04s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.12
+ botocore==1.43.12
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt=... [truncated, 41865 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes. This resulted in repeated OSError exceptions ("AF_UNIX path length cannot exceed 107 bytes") for all models, causing the fit process to abort with RuntimeError: "No models were trained successfully during fit()".

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output, especially the scratch and output directories, so that the full UNIX socket paths remain under 107 bytes. Move your working directory (e.g., `/sc-scratch/sc-scratch-ikim-guidlight/be-fair/mlzero/29-basic_prompt/node_15/`) to a much shorter path (such as `/tmp/ag/`), update all relevant environment variables and script paths accordingly, and rerun the training.

Node 16 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-21 23:39:36
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	16
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinde
stage	debug
time_step	16
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	7
child_ids	[]

Error Message:

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
Using Python 3.11.15 environment at: 29-basic_prompt/node_16/conda_env
Resolved 239 packages in 3.48s
Downloading botocore (14.4MiB)
Downloaded botocore
Prepared 2 packages in 12.04s
Uninstalled 1 package in 1.81s
Installed 237 packages in 44.44s
+ absl-py==2.4.0
+ accelerate==1.13.0
+ adagio==0.2.6
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiohttp-cors==0.8.1
+ aiosignal==1.4.0
+ alembic==1.18.4
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ antlr4-python3-runtime==4.9.3
+ attrs==26.1.0
+ autogluon==1.5.1b20260521
+ autogluon-common==1.5.1b20260521
+ autogluon-core==1.5.1b20260521
+ autogluon-features==1.5.1b20260521
+ autogluon-multimodal==1.5.1b20260521
+ autogluon-tabular==1.5.1b20260521
+ autogluon-timeseries==1.5.1b20260521
+ beartype==0.22.9
+ beautifulsoup4==4.14.3
+ blis==1.3.3
+ boto3==1.43.13
+ botocore==1.43.13
+ catalogue==2.0.10
+ catboost==1.2.10
+ certifi==2026.5.20
+ cffi==2.0.0
+ chardet==7.4.3
+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2
+ click==8.4.0
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.8... [truncated, 41899 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes, resulting in OSError: "validate_socket_filename failed: AF_UNIX path length cannot exceed 107 bytes". This caused all model training attempts to fail, leading to a RuntimeError indicating that no models were successfully trained.

SUGGESTED_FIX:
Shorten the directory paths used for temporary files and model output by moving your working, scratch, and output directories to locations with much shorter absolute paths (e.g., /tmp/ag_run or /scratch/ag_run). Set the TMPDIR environment variable to this shorter path before running your script, and ensure all AutoGluon and Ray-related temporary files are created there. This will prevent UNIX socket path length overflows and allow model training to proceed.