

Node Structure Visualization


Click on any node in the tree to jump to its detailed information.

Total Nodes: 19 | Best Validation Score: 0.9190 (Node 15)

Node Tree:

Table of Contents:

Node ID	Stage	Status	Tool
>-1			
>0			autogluon.multimodal
>1			machine learning
>2			autogluon.tabular
>3			autogluon.multimodal
>4			machine learning
>5			autogluon.multimodal
>6			machine learning
>7			autogluon.tabular
>8			autogluon.multimodal
>9			autogluon.multimodal
>10			machine learning
>11			autogluon.tabular
>12			autogluon.multimodal
>13			machine learning
>14			autogluon.tabular
>15			machine learning
>16			autogluon.tabular
>17			autogluon.tabular



Status	Color
Success	Green
Failure	Red
Neutral	Blue

Node Details:

Node -1 (root)

Property	Value
uct_value	0.5833
ctime	2026-05-22 16:55:50
debug_attempts	0
depth	0
execution_time	0.0
expected_child_count	0
failure_visits	17
id	-1
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
num_children	4
prev_tutorial_prompt	None
stage	root
time_step	-1
tool_used	
tools_available	[]
unvalidated_visits	0
validated_reward	0.919
validated_visits	1
validation_score	None
visits	18
parent_id	None
child_ids	[0, 1, 2, 15]

Node 0 (evolve)

Property	Value
uct_value	0.6481
ctime	2026-05-22 16:56:49
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.3
exploration	0.9968262051089475
failure_penalty	-0.3
failure_visits	6
id	0
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	3
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	0
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	6
parent_id	-1
child_ids	[3, 5]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_0/conda_env

Resolved 235 packages in 2.47s

Uninstalled 1 package in 2.08s

Installed 233 packages in 2m 19s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1
- + hf-xet==1.5.0
- + httpx==1.0.dev3
- + huggingface-hub==0.36.2
- + hyperopt==0.2.7

+ idna==3.16
... [truncated, 6620 chars total]

Error Analysis:

ERROR_SUMMARY: The error occurs because the `MultiModalPredictor.fit()` method in the installed version of AutoGluon MultiModal does not accept a `verbose` keyword argument. Passing `verbose=True` to `fit()` causes a `TypeError`, as this parameter is not defined in the method's signature.

SUGGESTED_FIX:

Remove the `verbose` argument from the call to `predictor.fit()`. The method will use its default logging behavior, and if you need to control verbosity, refer to the official documentation for supported logging or verbosity options.

Node 1 (evolve)

Property	Value
uct_value	0.7751
ctime	2026-05-22 17:07:48
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.3
exploration	0.7804660679599289
failure_penalty	-0.3
failure_visits	5
id	1
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	3
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	1
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	5
parent_id	-1
child_ids	[4, 6]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

warning: Requirements file `~/home/anri21/.conda/envs/mlzero\_env/lib/python3.11/site-packages/autogluon/assistant/tools\_registry/machine learning/requirements.txt` does not contain any dependencies

Using Python 3.11.15 environment at: 54-addition_3/node_1/conda_env

Resolved 79 packages in 770ms

Uninstalled 1 package in 2.37s

Installed 78 packages in 1m 00s

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiosignal==1.4.0

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ anyio==4.13.0

+ attrs==26.1.0

+ certifi==2026.5.20

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ click==8.4.1

+ cuda-bindings==13.2.0

+ cuda-pathfinder==1.5.4

+ cuda-toolkit==13.0.2

+ datasets==4.8.5

+ dill==0.4.1

+ filelock==3.29.0

+ frozenlist==1.8.0

+ fsspec==2026.2.0

+ h11==0.16.0

+ hf-xet==1.5.0

+ httpcore==1.0.9

+ httpx==0.28.1

+ huggingface-hub==1.16.1

+ idna==3.16

+ jinja2==3.1.6

+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocessing==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.14.0a1
+ pydantic-core==2.47.0
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2
+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 15086 chars total]

Error Analysis:

ERROR_SUMMARY: The process was terminated after exceeding the 1-hour wall-clock time limit, likely due to slow data loading and/or training. Repeated `OSError: [Errno 16] Device or resource busy` errors occurred during temporary file cleanup, which are typical on NFS filesystems when files are still open by other processes (often caused by too many DataLoader workers). Additionally, the script requested 16 DataLoader workers, while the system recommends a maximum of 2, which can cause slowness or freezing.

SUGGESTED_FIX:

Reduce the number of DataLoader workers to 2 or fewer by setting ``num_workers=2`` in all DataLoader instances to avoid resource contention and slowness. Monitor and limit batch size if memory is constrained, and consider reducing the number of training epochs or using a smaller model to ensure the script completes within the 1-hour time limit. After making these changes, rerun the script and check for completion and output correctness.

Node 2 (evolve)

Property	Value
uct_value	0.6481
ctime	2026-05-22 18:10:26
debug_attempts	2
depth	1
execution_time	0.0
expected_child_count	0
exploitation	-0.3
exploration	0.53219899601806
failure_penalty	-0.3
failure_visits	6
id	2
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	3
num_children	2
prev_tutorial_prompt	None
stage	evolve
time_step	2
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None

visits	6
parent_id	-1
child_ids	[11, 7]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_2/conda_env

Resolved 239 packages in 1.27s

Uninstalled 1 package in 1.95s

Installed 237 packages in 1m 55s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + daal==2025.9.0
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1
- + hf-xet==1.5.0
- + httpx==1.0.dev3
- + huggingface-hub==0.36.2

+ hyperopt==0.... [truncated, 6971 chars total]

Error Analysis:

ERROR_SUMMARY: The error occurs because the code attempts to select the columns ['skin_tone', 'alternative_skin_tone', 'expert_opinion'] from the test DataFrame, but these columns do not exist in test.csv. This results in a KeyError when executing `test_ag = test[feature_cols].copy()`.

SUGGESTED_FIX:

Update the code so that, when preparing the test set for prediction, you only select columns that are actually present in test.csv (i.e., only 'image_name'). Do not attempt to select or use columns like 'skin_tone', 'alternative_skin_tone', or 'expert_opinion' for the test set, as they are not available; adjust `test_ag` to use only the intersection of feature columns and test columns.

Node 3 (debug)

Property	Value
uct_value	0.9261
ctime	2026-05-22 18:17:34
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.2684446626418298
failure_penalty	-0.0
failure_visits	3
id	3
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	2
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementation
stage	debug
time_step	3
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	3
parent_id	0
child_ids	[8, 12]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_3/conda_env

Resolved 235 packages in 1.15s

Uninstalled 1 package in 1.79s

Installed 233 packages in 1m 38s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

- + catboost==1.2.10
- + certifi==2026.5.20
- + cffi==2.0.0
- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1
- + hf-xet==1.5.0

```
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.16
... [truncated, 9898 chars total]
```

Error Analysis:

ERROR_SUMMARY: The error occurs because the value passed to the focal loss alpha parameter (`optim.focal_loss.alpha``) is not in the expected format. Specifically, AutoGluon expects `alpha`` to be a list or numpy array of floats (one per class), but a string or an object with an incorrect shape/type was provided, resulting in the `ValueError: "too many dimensions 'str'"` when calling `torch.tensor(alpha, dtype=torch.float32)``.

SUGGESTED_FIX:

Check the value and type of `class_weights`` before passing it to the hyperparameters. Ensure that `class_weights`` is a Python list or numpy array of floats (e.g., `[0.5, 0.5]`` for binary classification), not a string or any other type. Print or log `class_weights`` right before model training to verify its structure, and confirm that it matches the number of classes and contains only numeric values.

Node 4 (debug)

Property	Value
uct_value	1.2684
ctime	2026-05-22 18:26:29
debug_attempts	2
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.2684446626418298
failure_penalty	-0.0
failure_visits	2
id	4
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have them
stage	debug
time_step	4
tool_used	machine learning

tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	2
parent_id	1
child_ids	[10]

Error Message:

```

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `~/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 54-addition_3/node_4/conda_env
Resolved 79 packages in 382ms
Uninstalled 1 package in 1.79s
Installed 78 packages in 41.08s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.1
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0

```

- + frozenlist==1.8.0
- + fsspec==2026.2.0
- + h11==0.16.0
- + hf-xet==1.5.0
- + httpcore==1.0.9
- + httpx==0.28.1
- + huggingface-hub==1.16.1
- + idna==3.16
- + jinja2==3.1.6
- + markdown-it-py==4.2.0
- + markupsafe==3.0.3
- + mdurl==0.1.2
- + mpmath==1.3.0
- + multidict==6.7.1
- + multiprocessing==0.70.19
- + networkx==3.6.1
- + numpy==2.4.6
- + nvidia-cublas==13.1.1.3
- + nvidia-cuda-cupti==13.0.85
- + nvidia-cuda-nvrtc==13.0.88
- + nvidia-cuda-runtime==13.0.96
- + nvidia-cudnn-cu13==9.20.0.48
- + nvidia-cufft==12.0.0.61
- + nvidia-cufile==1.15.1.6
- + nvidia-curand==10.4.0.35
- + nvidia-cusolver==12.0.4.66
- + nvidia-cusparselt-cu13==0.8.1
- + nvidia-cusparse==12.6.3.3
- + nvidia-nccl-cu13==2.29.7
- + nvidia-nvjitlink==13.0.88
- + nvidia-nvshmem-cu13==3.4.5
- + nvidia-nvtx==13.0.85
- + opencv-python==4.13.0.92
- + pandas==3.0.3
- + pillow==12.2.0
- + propcache==0.5.2
- + pyarrow==24.0.0
- + pydantic==2.14.0a1
- + pydantic-core==2.47.0
- + pygments==2.20.0
- + python-calamine==0.6.2
- + python-dateutil==2.9.0.post0
- + pytz==2026.2
- + pyyaml==6.0.3
- + requests==2.34.2

+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 5757 chars total]

Error Analysis:

ERROR_SUMMARY: The main error is that the process exceeded the 1-hour (3600 seconds) wall-clock time limit, likely due to slow data loading or model training. Additionally, there are repeated OSError warnings related to NFS resource contention (".nfs... Device or resource busy") during temporary directory cleanup, and a warning that the installed NVIDIA GPU driver is too old for the requested CUDA version, which may prevent GPU acceleration or cause fallback to CPU.

SUGGESTED_FIX:

Reduce the training time by decreasing the number of epochs, using a smaller model (e.g., a lighter EfficientNet variant or ResNet18), and/or lowering the batch size if I/O is a bottleneck. Investigate and resolve NFS file locking issues by ensuring no processes are holding onto deleted files, and consider using local disk for temporary files if possible. Update the NVIDIA GPU driver to a version compatible with your CUDA/PyTorch stack to enable GPU acceleration and avoid fallback to CPU, which significantly slows down training.

Node 5 (debug)

Property	Value
uct_value	1.3384
ctime	2026-05-22 19:29:18
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.664857771836881
failure_penalty	-0.0
failure_visits	2
id	5
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementation
stage	debug
time_step	5
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	2
parent_id	0
child_ids	[9]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_5/conda_env

Resolved 235 packages in 3.27s

Uninstalled 1 package in 2.19s

Installed 233 packages in 2m 46s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

- + catboost==1.2.10
- + certifi==2026.5.20
- + cffi==2.0.0
- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1
- + hf-xet==1.5.0

```
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.16
... [truncated, 11762 chars total]
```

Error Analysis:

ERROR_SUMMARY: The training process was terminated by a SIGTERM signal before completion, likely due to exceeding the wall-clock time limit (1 hour) imposed by the execution environment. This caused the AutoGluon MultiModal `.fit()` call to be interrupted, resulting in no model output or predictions.

SUGGESTED_FIX:

Reduce the total training time by lowering `optim.max_epochs` (e.g., to 3–5), decreasing model size (e.g., use a smaller backbone via `"model.timm_image.checkpoint_name"`), or reducing dataset size (e.g., subsample training data). Also, ensure that data loading and preprocessing are efficient, and set a stricter `time_limit` in `predictor.fit()` (e.g., 1000–1200 seconds) to leave time for inference and saving outputs before the hard timeout.

Node 6 (debug)

Property	Value
uct_value	1.2684
ctime	2026-05-22 19:41:51
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.664857771836881
failure_penalty	-0.0
failure_visits	2
id	6
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have them
stage	debug
time_step	6
tool_used	machine learning

tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	2
parent_id	1
child_ids	[13]

Error Message:

```

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `~/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 54-addition_3/node_6/conda_env
Resolved 79 packages in 622ms
Uninstalled 1 package in 1.99s
Installed 78 packages in 41.77s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.1
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0

```

- + frozenlist==1.8.0
- + fsspec==2026.2.0
- + h11==0.16.0
- + hf-xet==1.5.0
- + httpcore==1.0.9
- + httpx==0.28.1
- + huggingface-hub==1.16.1
- + idna==3.16
- + jinja2==3.1.6
- + markdown-it-py==4.2.0
- + markupsafe==3.0.3
- + mdurl==0.1.2
- + mpmath==1.3.0
- + multidict==6.7.1
- + multiprocessing==0.70.19
- + networkx==3.6.1
- + numpy==2.4.6
- + nvidia-cublas==13.1.1.3
- + nvidia-cuda-cupti==13.0.85
- + nvidia-cuda-nvrtc==13.0.88
- + nvidia-cuda-runtime==13.0.96
- + nvidia-cudnn-cu13==9.20.0.48
- + nvidia-cufft==12.0.0.61
- + nvidia-cufile==1.15.1.6
- + nvidia-curand==10.4.0.35
- + nvidia-cusolver==12.0.4.66
- + nvidia-cusparselt-cu13==0.8.1
- + nvidia-cusparse==12.6.3.3
- + nvidia-nccl-cu13==2.29.7
- + nvidia-nvjitlink==13.0.88
- + nvidia-nvshmem-cu13==3.4.5
- + nvidia-nvtx==13.0.85
- + opencv-python==4.13.0.92
- + pandas==3.0.3
- + pillow==12.2.0
- + propcache==0.5.2
- + pyarrow==24.0.0
- + pydantic==2.14.0a1
- + pydantic-core==2.47.0
- + pygments==2.20.0
- + python-calamine==0.6.2
- + python-dateutil==2.9.0.post0
- + pytz==2026.2
- + pyyaml==6.0.3
- + requests==2.34.2

+ rich==15.0.0
- setuptools==82.0.1
+ s... [truncated, 5363 chars total]

Error Analysis:

ERROR_SUMMARY: The main error is that the process was terminated after exceeding the 1-hour wall-clock time limit, meaning the script did not finish training and inference within the allowed 3600 seconds. Additional warnings indicate that the installed NVIDIA GPU driver is too old for the requested CUDA version (causing PyTorch to fall back to CPU), and there are "Device or resource busy" errors from NFS when trying to delete temporary files, but these are secondary to the timeout.

SUGGESTED_FIX:

Reduce the total runtime by further decreasing the number of training epochs (e.g., to 1–2), using a smaller model (e.g., resnet18 instead of efficientnet_b0), and/or reducing the training dataset size (e.g., subsample for quick prototyping). Also, check your environment's GPU driver compatibility and update the NVIDIA driver if possible to enable GPU acceleration, which will significantly speed up training and inference.

Node 7 (debug)

Property	Value
uct_value	0.9261
ctime	2026-05-22 20:44:18
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.2684446626418298
failure_penalty	-0.0
failure_visits	3
id	7
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	2
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinder dataset.
stage	debug
time_step	7
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	3
parent_id	2
child_ids	[16, 17]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_7/conda_env

Resolved 239 packages in 1.51s

Uninstalled 1 package in 1.65s

Installed 237 packages in 2m 43s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

- + catboost==1.2.10
- + certifi==2026.5.20
- + cffi==2.0.0
- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + daal==2025.9.0
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1

+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.... [truncated, 42698 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes. This resulted in repeated OSError exceptions (`validate\_socket\_filename failed: AF\_UNIX path length cannot exceed 107 bytes`), causing all model training to be skipped and no models to be fit.

SUGGESTED_FIX:

Shorten the directory paths used for temporary files and model output—specifically, move your working/output directories to a location with a much shorter absolute path (e.g., `/tmp/ag\_run` or `/scratch/ag\_run`). Ensure that both the `OUTPUT\_DIR` and any temp directories (including those used by Ray and AutoGluon) are as short as possible, then re-run your script. If you cannot change the base directory, set the `TMPDIR`, `TEMP`, and `TMP` environment variables to a short path before running Python, so that Ray and other libraries use a shorter temp path.

Node 8 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-22 20:53:14
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	8
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementation
stage	debug
time_step	8
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	3
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_8/conda_env

Resolved 235 packages in 1.55s

Uninstalled 1 package in 2.46s

Installed 233 packages in 1m 51s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2

+ click==8.4.1

+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.16
... [truncated, 11643 chars total]

Error Analysis:

ERROR_SUMMARY: The training process was terminated by a SIGTERM (signal 15), which is typically sent by the SLURM workload manager when the job exceeds its wall-clock time limit (1 hour in this case). As a result, the model did not complete training, and no predictions or validation scores were produced.

SUGGESTED_FIX:

Reduce the total training time by lowering the number of epochs (e.g., set `optim.max_epochs` to 2 or 3), decrease the model size (use a lighter backbone such as `mobilenetv3_large_100`), or further restrict the training time via the `time_limit` parameter (e.g., `time_limit=600`). Additionally, ensure that data loading and preprocessing are efficient, and consider using fewer GPUs or workers if initialization overhead is high.

Node 9 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-22 21:03:55
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	9
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementation
stage	debug
time_step	9
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	5
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_9/conda_env

Resolved 235 packages in 2.17s

Uninstalled 1 package in 1.98s

Installed 233 packages in 1m 41s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2

+ click==8.4.1

+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0.2.7
+ idna==3.16
... [truncated, 11645 chars total]

Error Analysis:

ERROR_SUMMARY: The training process was terminated by a SIGTERM signal, which is typically sent by the job scheduler (SLURM) when the script exceeds its allocated wall-clock time or due to resource management policies. As a result, the model did not complete training, and no predictions or validation metrics were produced.

SUGGESTED_FIX:

Reduce the total runtime by further lowering the number of training epochs (e.g., set `optim.max\_epochs` to 1 or 2), decrease model complexity (use a smaller/faster backbone), and/or reduce dataset size (use a smaller training subset or lower image resolution). Additionally, ensure that data loading and preprocessing are efficient, and consider profiling the script to identify and optimize any slow steps.

Node 10 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-22 21:14:26
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	10
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have them
stage	debug
time_step	10
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0

validation_score	None
visits	1
parent_id	4
child_ids	[]

Error Message:

```

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `~/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 54-addition_3/node_10/conda_env
Resolved 79 packages in 376ms
Uninstalled 1 package in 1.80s
Installed 78 packages in 38.40s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.1
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0
+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.1
+ idna==3.16

```

```
+ jinja2==3.1.6
+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocessing==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.14.0a1
+ pydantic-core==2.47.0
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2
+ rich==15.0.0
- setuptools==82.0.1
+ ... [truncated, 7010 chars total]
```

Error Analysis:

ERROR_SUMMARY: The error "RuntimeError: mat1 and mat2 shapes cannot be multiplied (32x1280 and 960x1)" occurs because the classifier head of the MobileNetV3 model is initialized with an incorrect input feature size (in_features=960), while the actual output feature size from the backbone is 1280. This shape mismatch prevents the linear layer from performing matrix multiplication during the forward pass.

SUGGESTED_FIX: Check the value assigned to `in_features` when constructing the classifier head in your `MobileNetV3Classifier`. Ensure that `in_features` matches the output feature dimension of the backbone (which is 1280 for "mobilenetv3_large_100") by referencing `self.backbone.num_features` or the correct attribute from the `timm` model. Re-initialize the head layer with `in_features=1280` to resolve the mismatch.

Node 11 (debug)

Property	Value
uct_value	1.3384
ctime	2026-05-22 21:19:54
debug_attempts	3
depth	2
execution_time	0.0
expected_child_count	0
exploitation	0.0
exploration	1.482079962591042
failure_penalty	-0.0
failure_visits	2
id	11
is_debug_successful	✗
is_leaf	✗
is_successful	✗
is_terminal	✓
normalized_failure_visit	0
num_children	1
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates using AutoGluon Multimodal for multimodal classification on the PetFinder dataset.
stage	debug
time_step	11
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_contribution	0.0
validated_reward	0.0

validated_visits	0
validation_score	None
visits	2
parent_id	2
child_ids	[14]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_11/conda_env

Resolved 239 packages in 1.32s

Uninstalled 1 package in 2.04s

Installed 237 packages in 1m 43s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.13

+ botocore==1.43.13

+ catalogue==2.0.10

- + catboost==1.2.10
- + certifi==2026.5.20
- + cffi==2.0.0
- + chardet==7.4.3
- + charset-normalizer==3.4.7
- + chronos-forecasting==2.2.2
- + click==8.4.1
- + cloudpathlib==0.24.0
- + cloudpickle==3.1.2
- + colorama==0.4.6
- + colorful==0.6.0a1
- + colorlog==6.10.1
- + confection==1.3.3
- + contourpy==1.3.3
- + coreforecast==0.0.16
- + cryptography==48.0.0
- + cycler==0.12.1
- + cymem==2.0.13
- + daal==2025.9.0
- + datasets==4.0.0
- + defusedxml==0.7.1
- + dill==0.3.8
- + distlib==0.4.0
- + einops==0.8.2
- + einx==0.4.3
- + evaluate==0.4.6
- + fastai==2.8.5
- + fastcore==1.8.18
- + fastdownload==0.0.7
- + fastprogress==1.0.5
- + fasttransform==0.0.2
- + filelock==3.29.0
- + fonttools==4.63.0
- + frozendict==2.4.7
- + frozenlist==1.8.0
- + fsspec==2025.3.0
- + fugue==0.9.7
- + future==1.0.0
- + gdown==6.0.0
- + gluonts==0.16.2
- + google-api-core==2.30.3
- + google-auth==2.53.0
- + googleapis-common-protos==1.75.0
- + graphviz==0.21
- + grpcio==1.81.0rc1

```
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0... [truncated, 42962 chars total]
```

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the UNIX socket path used by Ray (a backend for parallel processing) exceeded the maximum allowed length (107 bytes) for AF_UNIX sockets. This caused OSError exceptions during model training, resulting in all models being skipped and a final RuntimeError indicating that no models were trained. The root cause is the excessive length of the temporary directory path, likely due to deep directory nesting in your scratch space.

SUGGESTED_FIX:

Shorten the directory path used for temporary files and model output by moving your working directory (including data, output, and conda environment) to a location with a much shorter absolute path (e.g., /tmp/youruser/ or /scratch/youruser/). Ensure that all relevant environment variables (such as TMPDIR, TEMP, and TMP) point to this shorter path before running your script, so that Ray and AutoGluon use it for temporary files and socket creation. Then, rerun your training script.

Node 12 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-22 21:27:20
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	12
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial demonstrates image classification using AutoGluon MultiModal, covering implementation
stage	debug
time_step	12
tool_used	autogluon.multimodal
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	3
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_12/conda_env

Resolved 235 packages in 2.99s

Downloading botocore (14.4MiB)

Downloaded botocore

Prepared 2 packages in 12.05s

Uninstalled 1 package in 1.84s

Installed 233 packages in 1m 50s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.14

+ botocore==1.43.14

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2
+ click==8.4.1
+ cloudpathlib==0.24.0
+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.... [truncated, 11733 chars total]

Error Analysis:

ERROR_SUMMARY: The training process was terminated by a SIGTERM signal, which is typically sent by the job scheduler (SLURM) due to exceeding the wall-clock time limit or resource constraints. As a result, the model did not complete training, and no predictions or validation metrics were produced.

SUGGESTED_FIX:

Reduce the total training time by lowering the number of epochs, using a smaller model backbone, or decreasing dataset size (e.g., subsample for development). Additionally, ensure that the combined time for data loading, training, and inference stays well below the 1-hour wall-clock limit imposed by your environment. Monitor actual runtime in logs and adjust `optim.max_epochs``, `time_limit``, and batch size accordingly to guarantee completion within the allowed time.

Node 13 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-22 21:38:08
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	13
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: Intro to Use Machine Learning or Deep Learning algorithms # Condensed: Intro to Use Machine Learning or Deep Learning algorithms Summary: You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you # Intro to Use Machine Learning or Deep Learning algorithms You can use any machine learning or deep learning algorithms to complete the tasks. Make sure you have them
stage	debug
time_step	13
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0

validation_score	None
visits	1
parent_id	6
child_ids	[]

Error Message:

```

stderr:
==> WARNING: A newer version of conda exists. <==
current version: 25.11.0
latest version: 26.5.0
Please update conda by running
$ conda update -n base -c conda-forge conda
warning: Requirements file `~/home/anri21/.conda/envs/mlzero_env/lib/python3.11/site-packages/autogluon/assistant/tools_registry/machine learning/requirements.txt`
does not contain any dependencies
Using Python 3.11.15 environment at: 54-addition_3/node_13/conda_env
Resolved 79 packages in 345ms
Uninstalled 1 package in 1.64s
Installed 78 packages in 42.75s
+ aiohappyeyeballs==2.6.2
+ aiohttp==3.13.5
+ aiosignal==1.4.0
+ annotated-doc==0.0.4
+ annotated-types==0.7.0
+ anyio==4.13.0
+ attrs==26.1.0
+ certifi==2026.5.20
+ chardet==7.4.3
+ charset-normalizer==3.4.7
+ click==8.4.1
+ cuda-bindings==13.2.0
+ cuda-pathfinder==1.5.4
+ cuda-toolkit==13.0.2
+ datasets==4.8.5
+ dill==0.4.1
+ filelock==3.29.0
+ frozenlist==1.8.0
+ fsspec==2026.2.0
+ h11==0.16.0
+ hf-xet==1.5.0
+ httpcore==1.0.9
+ httpx==0.28.1
+ huggingface-hub==1.16.1
+ idna==3.16

```

```
+ jinja2==3.1.6
+ markdown-it-py==4.2.0
+ markupsafe==3.0.3
+ mdurl==0.1.2
+ mpmath==1.3.0
+ multidict==6.7.1
+ multiprocessing==0.70.19
+ networkx==3.6.1
+ numpy==2.4.6
+ nvidia-cublas==13.1.1.3
+ nvidia-cuda-cupti==13.0.85
+ nvidia-cuda-nvrtc==13.0.88
+ nvidia-cuda-runtime==13.0.96
+ nvidia-cudnn-cu13==9.20.0.48
+ nvidia-cufft==12.0.0.61
+ nvidia-cufile==1.15.1.6
+ nvidia-curand==10.4.0.35
+ nvidia-cusolver==12.0.4.66
+ nvidia-cusparselt-cu13==0.8.1
+ nvidia-nccl-cu13==2.29.7
+ nvidia-nvjitlink==13.0.88
+ nvidia-nvshmem-cu13==3.4.5
+ nvidia-nvtx==13.0.85
+ opencv-python==4.13.0.92
+ pandas==3.0.3
+ pillow==12.2.0
+ propcache==0.5.2
+ pyarrow==24.0.0
+ pydantic==2.14.0a1
+ pydantic-core==2.47.0
+ pygments==2.20.0
+ python-calamine==0.6.2
+ python-dateutil==2.9.0.post0
+ pytz==2026.2
+ pyyaml==6.0.3
+ requests==2.34.2
+ rich==15.0.0
- setuptools==82.0.1
+ ... [truncated, 6886 chars total]
```

Error Analysis:

ERROR_SUMMARY: The error "RuntimeError: mat1 and mat2 shapes cannot be multiplied (32x1280 and 960x1)" occurs because the output feature size of the MobileNetV3 backbone (1280) does not match the expected input size (960) of the classifier head. This mismatch is due to using `self.backbone.num_features` (which is 960 for some models, but 1280 for `mobilenetv3_large_100` in `timm`), leading to an incorrectly sized linear layer.

SUGGESTED_FIX:

Check the actual output feature dimension of the backbone by passing a dummy input through `self.backbone` and use that value to define the classifier head's input size. Replace `self.backbone.num_features` with the dynamically determined feature size (e.g., by running a forward pass with a dummy tensor of shape `[1, 3, 224, 224]`) to ensure the Linear layer matches the backbone output.

Node 14 (debug)

Property	Value
uct_value	1.1772
ctime	2026-05-22 21:43:24
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	14
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial explains GPU acceleration in AutoGluon, covering basic GPU allocation with `num_gpus`
stage	debug
time_step	14
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	11
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_14/conda_env

Resolved 239 packages in 1.39s

Uninstalled 1 package in 1.65s

Installed 237 packages in 1m 37s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.14

+ botocore==1.43.14

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2

+ click==8.4.1

+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0... [truncated, 42732 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create UNIX domain sockets with file paths exceeding the system limit of 107 bytes. This resulted in repeated OSError exceptions (`validate_socket_filename failed: AF_UNIX path length cannot exceed

107 bytes`), causing all model training attempts to be skipped and ultimately raising a `RuntimeError` due to no models being fit.

SUGGESTED_FIX:

Shorten the directory paths used for temporary files and model output by moving your working directories (including scratch, tmp, and output folders) to a location with a much shorter absolute path (e.g., ``tmp/your_project``). Update all relevant path variables in your scripts and environment to use these shorter paths, then rerun the training. This will ensure that Ray's socket file paths remain within the UNIX limit and allow model training to proceed.

Node 15 (debug)

Property	Value
uct_value	3.4040
ctime	2026-05-22 21:50:04
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	0
id	15
is_debug_successful	✗
is_leaf	✓
is_successful	✓
is_terminal	✓
num_children	0
prev_tutorial_prompt	None
stage	debug
time_step	15
tool_used	machine learning
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.919
validated_visits	1
validation_score	0.9190
visits	1
parent_id	-1
child_ids	[]

Node 16 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-22 22:14:52
debug_attempts	2
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	16
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial explains GPU acceleration in AutoGluon, covering basic GPU allocation with `num_gpus`
stage	debug
time_step	16
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	7
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_16/conda_env

Resolved 239 packages in 2.91s

Uninstalled 1 package in 1.60s

Installed 237 packages in 1m 34s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.14

+ botocore==1.43.14

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2

+ click==8.4.1

+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0... [truncated, 42837 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because Ray (used internally for parallelization) attempted to create UNIX domain socket files with paths longer than the system limit (107 bytes). This resulted in OSError: "AF_UNIX path length cannot exceed 107 bytes" for each model, causing all training attempts to be

skipped and ultimately raising `RuntimeError: "No models were trained successfully during fit()"`.

SUGGESTED_FIX:

Shorten the directory paths used for temporary files and model output—move your working, scratch, and output directories to a location with a much shorter absolute path (e.g., ``tmp/ag_run`` or ``scratch/ag_run``). Update all relevant environment variables and script paths accordingly, then rerun your pipeline to ensure Ray's socket files remain within the 107-byte limit.

Node 17 (debug)

Property	Value
uct_value	1.4821
ctime	2026-05-22 22:22:30
debug_attempts	3
depth	3
execution_time	0.0
expected_child_count	0
failure_visits	1
id	17
is_debug_successful	✗
is_leaf	✓
is_successful	✗
is_terminal	✓
num_children	0
prev_tutorial_prompt	### Condensed: ```python # Condensed: ```python Summary: This tutorial explains GPU acceleration in AutoGluon, covering basic GPU allocation with `num_gpus`
stage	debug
time_step	17
tool_used	autogluon.tabular
tools_available	['autogluon.multimodal', 'machine learning', 'autogluon.tabular']
unvalidated_visits	0
validated_reward	0.0
validated_visits	0
validation_score	None
visits	1
parent_id	7
child_ids	[]

Error Message:

stderr:

==> WARNING: A newer version of conda exists. <==

current version: 25.11.0

latest version: 26.5.0

Please update conda by running

\$ conda update -n base -c conda-forge conda

Using Python 3.11.15 environment at: 54-addition_3/node_17/conda_env

Resolved 239 packages in 1.20s

Uninstalled 1 package in 1.68s

Installed 237 packages in 1m 49s

+ absl-py==2.4.0

+ accelerate==1.13.0

+ adagio==0.2.6

+ aiohappyeyeballs==2.6.2

+ aiohttp==3.13.5

+ aiohttp-cors==0.8.1

+ aiosignal==1.4.0

+ alembic==1.18.4

+ annotated-doc==0.0.4

+ annotated-types==0.7.0

+ antlr4-python3-runtime==4.9.3

+ attrs==26.1.0

+ autogluon==1.5.1b20260522

+ autogluon-common==1.5.1b20260522

+ autogluon-core==1.5.1b20260522

+ autogluon-features==1.5.1b20260522

+ autogluon-multimodal==1.5.1b20260522

+ autogluon-tabular==1.5.1b20260522

+ autogluon-timeseries==1.5.1b20260522

+ beartype==0.22.9

+ beautifulsoup4==4.14.3

+ blis==1.3.3

+ boto3==1.43.14

+ botocore==1.43.14

+ catalogue==2.0.10

+ catboost==1.2.10

+ certifi==2026.5.20

+ cffi==2.0.0

+ chardet==7.4.3

+ charset-normalizer==3.4.7

+ chronos-forecasting==2.2.2

+ click==8.4.1

+ cloudpathlib==0.24.0

+ cloudpickle==3.1.2
+ colorama==0.4.6
+ colorful==0.6.0a1
+ colorlog==6.10.1
+ confection==1.3.3
+ contourpy==1.3.3
+ coreforecast==0.0.16
+ cryptography==48.0.0
+ cycler==0.12.1
+ cymem==2.0.13
+ daal==2025.9.0
+ datasets==4.0.0
+ defusedxml==0.7.1
+ dill==0.3.8
+ distlib==0.4.0
+ einops==0.8.2
+ einx==0.4.3
+ evaluate==0.4.6
+ fastai==2.8.5
+ fastcore==1.8.18
+ fastdownload==0.0.7
+ fastprogress==1.0.5
+ fasttransform==0.0.2
+ filelock==3.29.0
+ fonttools==4.63.0
+ frozendict==2.4.7
+ frozenlist==1.8.0
+ fsspec==2025.3.0
+ fugue==0.9.7
+ future==1.0.0
+ gdown==6.0.0
+ gluonts==0.16.2
+ google-api-core==2.30.3
+ google-auth==2.53.0
+ googleapis-common-protos==1.75.0
+ graphviz==0.21
+ grpcio==1.81.0rc1
+ hf-xet==1.5.0
+ httpx==1.0.dev3
+ huggingface-hub==0.36.2
+ hyperopt==0... [truncated, 42765 chars total]

Error Analysis:

ERROR_SUMMARY: AutoGluon failed to train any models because the Ray backend (used for parallel model training) attempted to create Unix domain sockets with file paths exceeding the system limit of 107 bytes. This resulted in repeated OSError exceptions (`validate_socket_filename failed: AF_UNIX path length cannot exceed

107 bytes`), causing all model training attempts to be skipped and ultimately raising a RuntimeError due to no models being fit.

SUGGESTED_FIX:

Shorten the directory paths used for temporary files and model output, especially the scratch and output directories, so that the full socket file paths generated by Ray remain under 107 bytes. Move your working directories (e.g., `/sc-scratch/sc-scratch-ikim-guidlight/...`) to a location with a much shorter absolute path (such as `/tmp/agwork`), update all relevant path variables in your scripts accordingly, and rerun the training.